

Project Portfolio by Sebastian Jost

- Short, visual summaries of personal projects
- Working time includes all aspects of a project: planning, implementation, documentation and writing of a paper or report (if applicable)
- Some linked GitHub repositories are private. Please request access if interested.
-

Solving General Twisty Puzzles with RL

Goal: Compute interpretable solutions on consumer hardware


Methods:

- Automatic action space augmentation
- Dense, intuitive rewards

Result:

- 100x faster training than previous RL-approaches
- AI solutions similar to human ones

Time: ~600 hours over 4 years



[Code on GitHub](#)

master's thesis

video explanation

[more demo videos](#)

Simulating real-world traffic sensor noise for more robust RL-Traffic Signal Control

Goal: Improve robustness of RL-TSC agents against noisy data

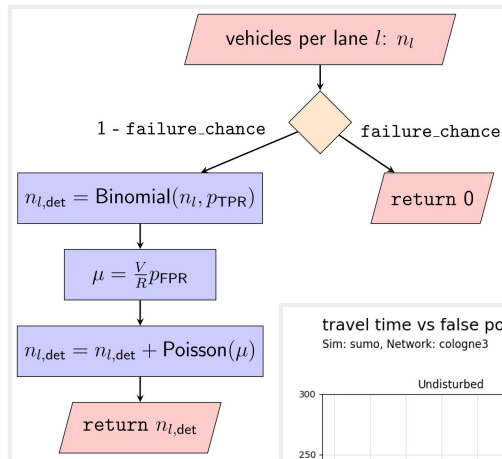
Methods:

- Simulate realistic sensor noise
- >600,000 simulated hours to test parameter settings and analyze their impact

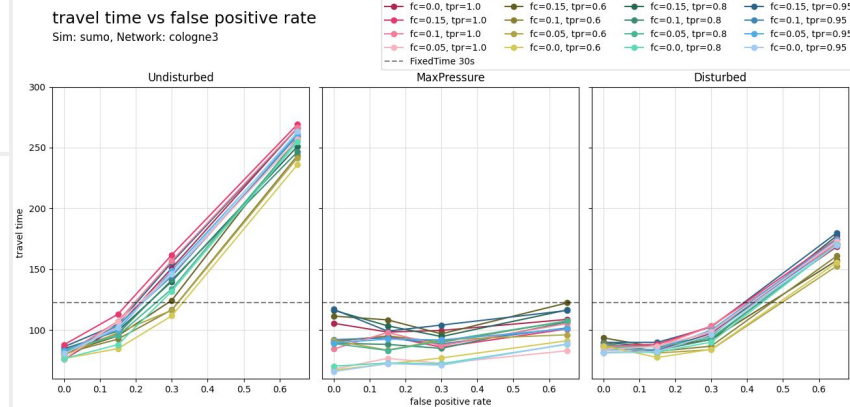
Result:

- Understanding impact of different noise types on tested RL and non-RL agents
- Actionable sensor type deployment recommendations

Time: ~200 hours incl. paper writing



Noisy calculation of detected vehicles from actual vehicles



[Code on GitHub](#)

publication planned

Custom Ticket to Ride Map Creation Tools

Goal: Efficiently create custom Ticket to Ride versions


Methods:

- Analyzed creation bottlenecks and processes
- Built tools to support creation at multiple stages

Result:

- Faster creation of variants
- High-quality custom games

Time: ~300 hours over 1 year



[Code on GitHub](#)

[result showcase](#)

Northgard Board Game with Web-UI

Goal:

- Recreate the board game Northgard digitally
- Write web-UI without prior JS/CSS/html knowledge

Methods:

- AI tools for writing frontend
- Python backend
- Custom tools for asset creation

Result:

- Prototype with responsive web-UI

Time: ~100 hours incl. asset creation



[Code on GitHub](#)

[demo video](#)

[game info](#)

[Assets on GitHub](#)

Simplified Saboteur Game with RL players

Goal:

- Investigate emergence of deceptive AI behaviour

Methods:

- LLM-accelerated coding

Result:

- RL-agents successfully play simplified game
- *unfinished*

Time: ~20 hours



[Code on GitHub](#)

[demo video](#)

Learning Methods in Humans, Animals and AI

Goal: Compare learning techniques in different organisms to inspire future research

Methods:

- Offered Workshop at [VANconference25](#)
- Years of following current research

Result:

- Overview helps predict industry & research direction

Time: <10 hours

Learning method	Animals	Humans	Computer (Machine Learning)
Evolution	genetic evolution	genetic evolution	genetic algorithms (e.g. NEAT)
Trial & Error	Conditioning	Conditioning	Reinforcement learning
Imitation	Orcas teaching hunting methods & language	Learning language	Supervised learning, Imitation learning
Abstract Instructions	Bee dance to describe location of food	Most of school / higher education, baking recipes	In-context learning in LLMs
Mental Simulation	Birds cracking nuts using cars, tool use	Dreams, pondering other's reactions	World models (e.g. DreamerV3)

[GitHub page](#)

Maze Generator & Solver using Q-Learning

Goal:

- Automatic maze generation
- Q-learning based solver


Methods:

- Customized random walk
- Q-learning
- tkinter GUI design

Result:

- Solver for small mazes
- Clear state-value visualization

Time: 3 days \approx 20 hours



Wizard Game with multiple AI players

Goal: Compare different AI techniques for creating wizard game agents


Methods:

- Rule-based AI, optimized via genetic evolution
- tkinter GUI design

Result:

- Major capability differences between AI agents

Time: ~60 hours



[Code on GitHub](#)